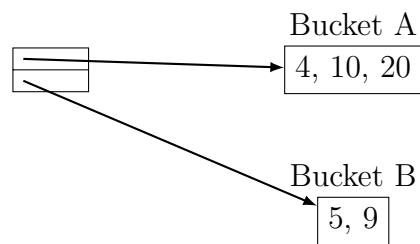
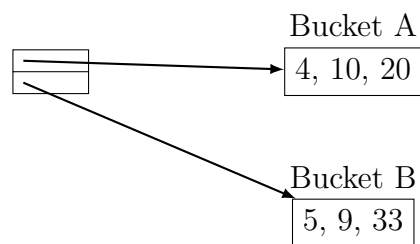


Problem 1

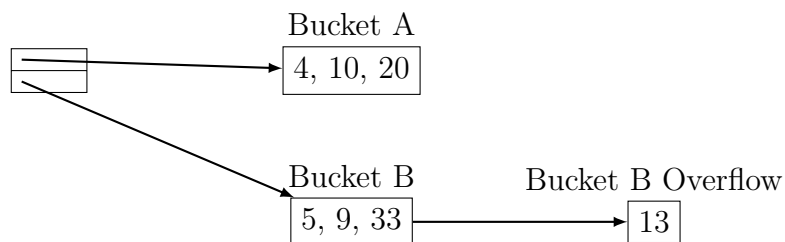
After inserting 20:




After inserting 33:



After inserting 13:



After inserting 14:



Problem 2

1

- (a) Non-blocking, because it can output tuples as it processes inputs, and just won't output them again if it has encountered the same tuple before.
- (b) Non-blocking. If column X is sorted, once it finds a different value of X it can output all tuples of the previous value because they make up an entire group.

- (c) Blocking, because it needs to find all elements in a specific group before it can start outputting them.
- (d) Blocking, because it needs to process all tuples in R before it can output the sorted list of tuples.
- (e) Non-blocking. Since the leaves of the B-tree are already sorted it can just output them in order as it reads them in.
- (f) Blocking, because it must first sort R and S, and then merge join them.
- (g) Non-blocking, it can output the resulting tuples as it reads in the input.

2

- (a) Can be done in one pass assuming the distinct tuples of R fit in 199 buffers.
- (b) Can be done in one pass as long as the biggest group can fit in 199 buffers.
- (c) Can be done in one pass as long as R can fit in 199 buffers.
- (d) Cannot be done in one pass. A two-pass external sort reads M blocks at a time, sorts them, and writes them to the disk as the first run, then merges the runs to produce a sorted output. The I/O cost will be $2 \times B(R)$ for the first pass and $B(R)$ for the second pass, for a total of $3B(R) = 3 \times 1000 = 3000$ I/Os.
- (e) Can be done in one pass, with 70 I/Os to read the index, plus 2000 I/Os to read and write the sorted relation R, for a total of 2070 I/Os.
- (f) Cannot be done in one pass. Phase one can sort R and S, while phase two merges and joins the sorted relations R and S. The I/O cost is $2 \times B(R) + 2 \times B(S)$ for phase one sorting, plus $B(R) + B(S)$ for the merge and join phase two, for a total cost of $3(B(R) + B(S)) = 3(1000 + 150) = 3450$ I/Os.
- (g) Can be done in one pass because S can fit in 199 buffers.

Problem 3

1. $Q = T(R1)/V(R1_a) = 400/50 = 8$
2. $Q = \left(\frac{T(R1)}{V(R1_a)} \right) \times 41 = 328$
3. $Q = 328 \times (1/50) = 7$
4. $Q = T(R1) \times T(R2) \div \max(V(R1_b), V(R2_b)) = 400 \times 500 \div \max(50, 40) = 4000$
5. $Q = 4000 \times T(R3) \div \max(V(R2_c), V(R3_c)) = 4000 \times 1000 \div 100 = 40000$
6. $Q = (328 \bowtie R2) \bowtie R3 = (328 \times 500 \div 50) \bowtie R3 = 3280 \bowtie R3 = 3280 \times 1000 \div 100 = 32800$

Problem 4

1.
 - a. R2, because it's smaller
 - b. $B(R2) + \frac{B(R2)}{M-1} \times B(R1) = 200 + \frac{200}{51} \times 1000 = 4122$ IOs
 - c. $B(R1) + \frac{B(R1)}{M-1} \times B(R2) = 1000 + \frac{1000}{51} \times 200 = 4922$ IOs
2.
 - a. $3(B(R1) + B(R2)) = 3(1000 + 200) = 3600$
 - b. $B(R1) + B(R2) \leq M^2$
 $1000 + 200 \leq M^2$
 $M \geq 34$
Minimum number of buffers necessary is 34
3.
 - a. $3(B(R1) + B(R2)) = 3(1000 + 200) = 3600$
 - b. $B(R2) \leq M^2$
 $200 \leq M^2$
 $M \geq 14$
Minimum number of buffers is 14
4.
 - a. Assuming R2 is clustered, R1.a index is in memory, and the index is clustered.
 $B(R2) + T(R2) \times (B(R1)/V(R1, a)) = 200 + 2000 \times (1000/500) = 4200$