Problem 1

(a) LRU Replacement Policy

Instruction:	1*	2*	3*	1	4*	5*	3	4	1	6	7	8*	9*	5	10
Frame 1	**1**	**1**	**1**	1	1	1	1	1	1	1	1	1	1	1	10
Frame 2		**2**	**2**	**2**	**2**	**2**	**2**	**2**	**2**	**2**	**2**	**2**	**2**	**2**	**2**
Frame 3			**3**	**3**	**3**	**3**	3	3	3	3	3	**8**	**8**	**8**	**8**
Frame 4					**4**	**4**	**4**	4	4	4	4	4	**9**	**9**	**9**
Frame 5						**5**	**5**	**5**	**5**	**5**	**5**	**5**	**5**	**5**	**5**
Frame 6										6	6	6	6	6	6
Frame 7											7	7	7	7	7

Note: pages in bold are pinned.

(b) Clock Replacement Policy

Note: a subscrpt p indicates that a page is pinned, and a subscript r indicates that the reference bit is set to 1.

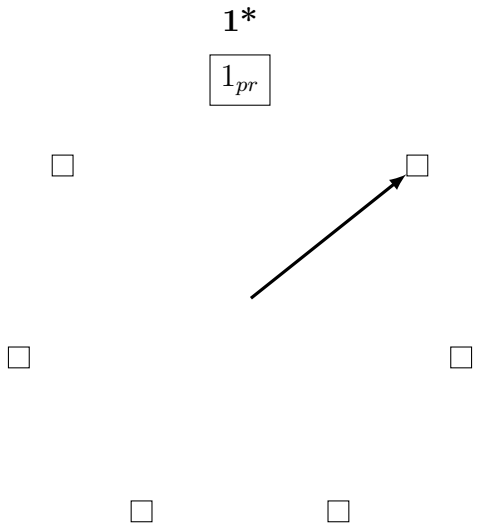

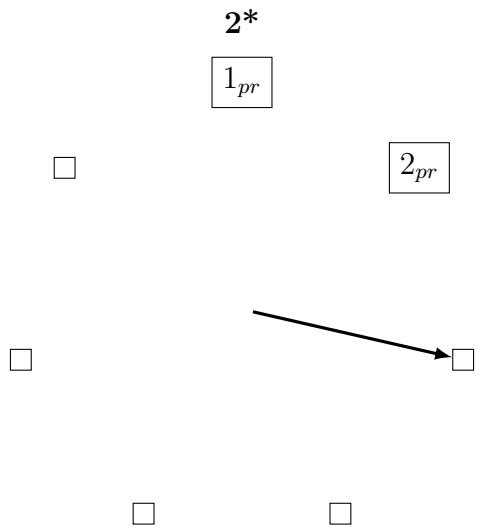
2*
1pr
2pr

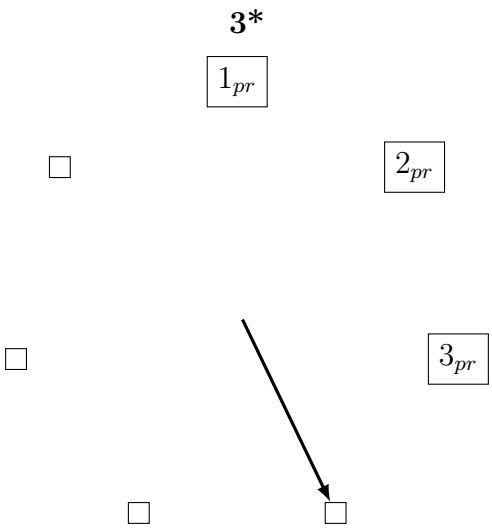
3*
1pr
2pr
3pr

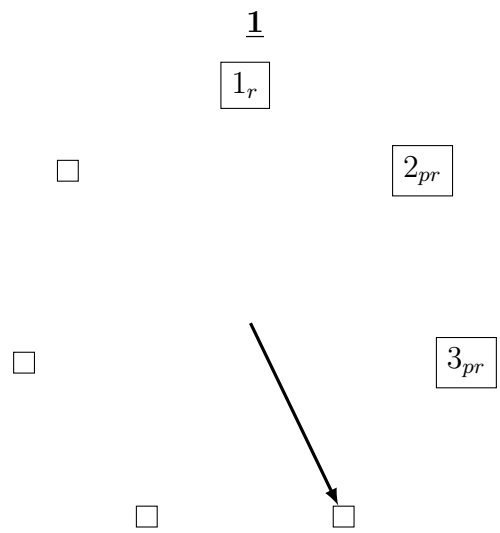
1
1r
2pr
3pr

4*

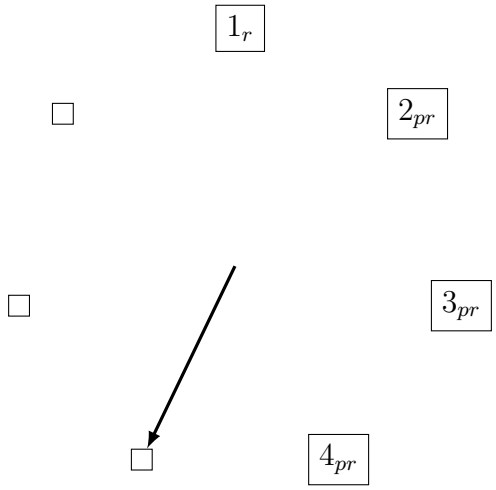

5*

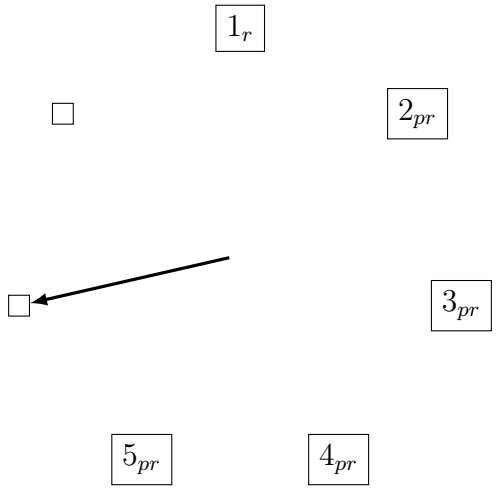

3

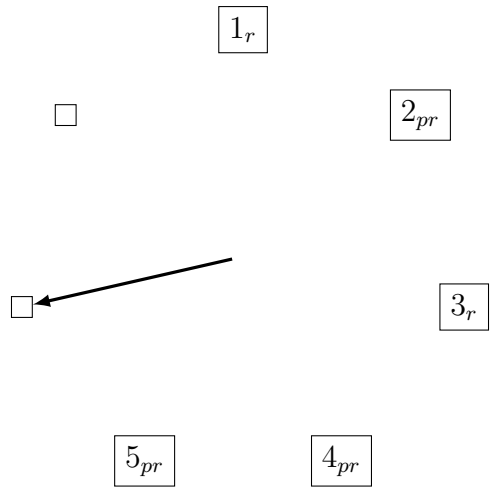

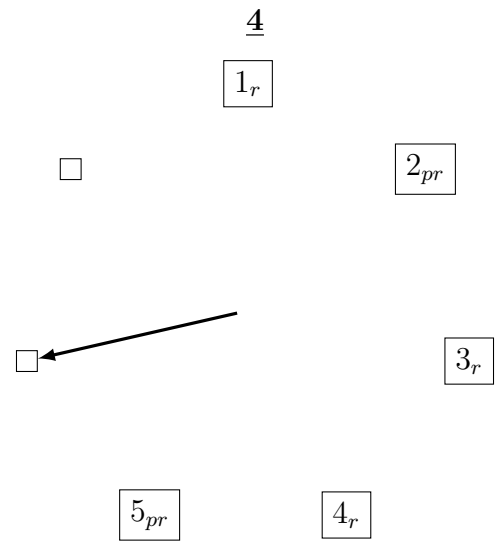
4
1r
2pr
3r
5pr
4r

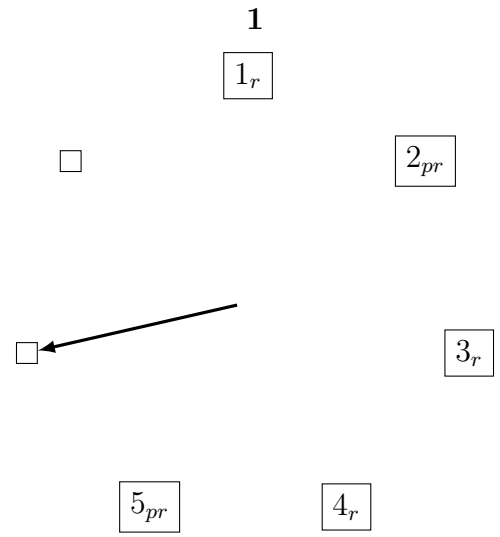
1
1r
2pr
3r
5pr
4r

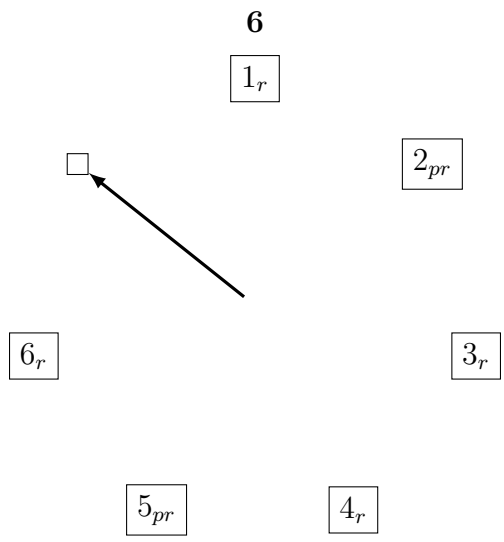
6
1r
2pr
6r
3r
5pr
4r

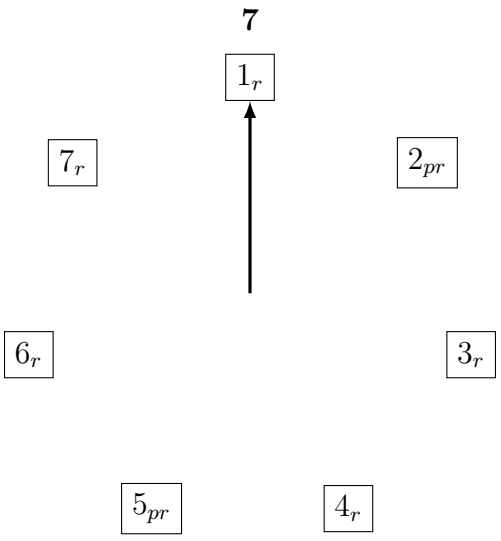
7
1_r
7_r
2_{pr}
6_r
3_r
5_{pr}
4_r

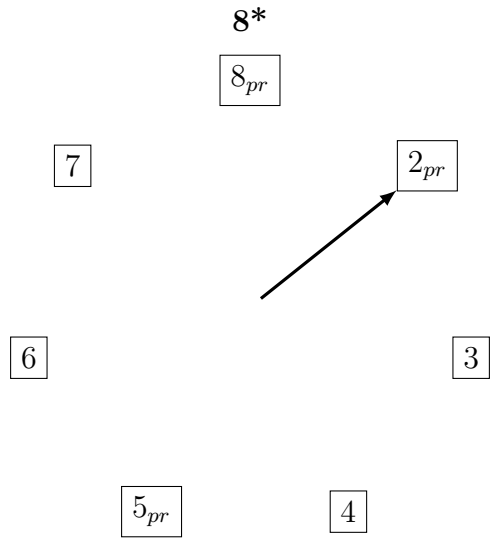
8*
8_{pr}
7
2_{pr}
6
3
5_{pr}
4

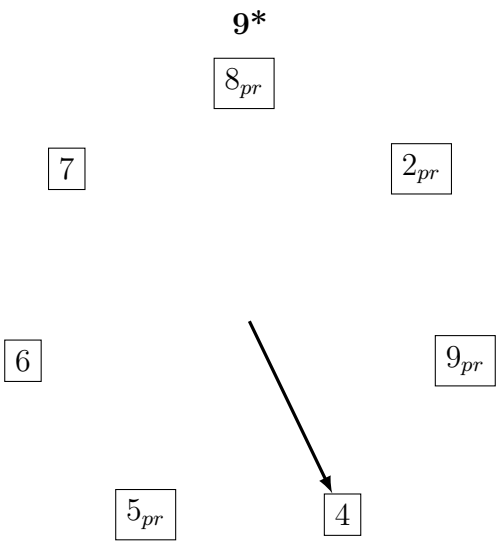
9*
8_{pr}
7
2_{pr}
6
9_{pr}
5_{pr}
4

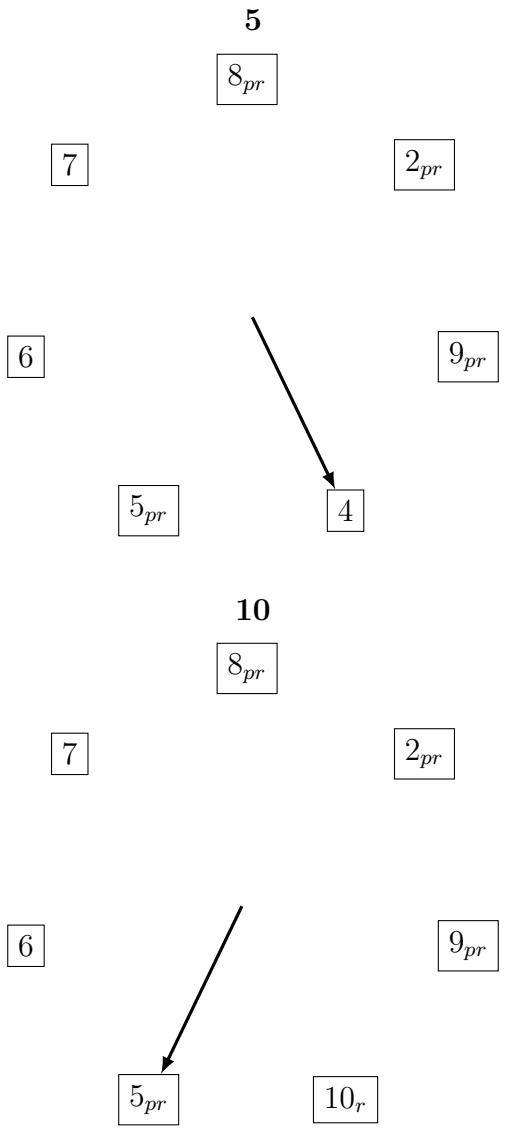

Problem 2

1. Yes. Each entry in the index would contain the value of K1 (20 bytes) and a pointer to the corresponding record (8 bytes). The entire data file contains 20 million records, so the index file would be 28 million bytes. Each block can contain $\lfloor 8,192/28 \rfloor = 292$ index records. The entire index would therefore be $\lceil 28,000,000/292 \rceil = 95,891$ blocks.

2. Yes. Each entry in the index would contain a pointer to the start of a block in the data file, and the value of K1 for the first record in that block. Since each block in the data file contains 20 records, you would need $1,000,000/20 = 50,000$ entries in the index, which would take up $\lceil 50,000/292 \rceil = 172$ blocks.